

# A path to Ethereum 2.0

(one guy's current understanding)

2018-07-06

Ben Edgington



# Background

- Scalability workshop, Taipei, 19-21 March 2018
- Client devs' and scaling workshop, Berlin, 29 June - 1 July 2018
- Lots of activity on [ethresear.ch](https://ethresear.ch) and [elsewhere](#)
- Ethereum [Core Devs calls](#)

## Disclaimer

None of this is set in stone; everything is subject to change. My understanding is likely to be faulty and/or incomplete. I am not party to any special knowledge: this is just an attempted synthesis of publicly available information.

**Update 2018-07-08:** I've added a (very) speculative delivery timeline expanding on some [remarks](#) from Justin Drake.

## My nomenclature

Ethereum 1.0 is the current Mainnet, which continues as a PoW chain for the foreseeable future.

Ethereum 2.0 is features likely to be delivered “soon”, possibly in sub-stages.

Ethereum 2.x is loosely features possibly to be implemented later. Maybe 3.0.

# Recent developments (last 3-4 weeks)

- Rather than upgrade the current Mainnet (Ethereum 1.0), innovation will happen on a separate Ethereum 2.0 blockchain (loosely coupled to Mainnet).
  - Technical reasons:
    - Capacity of Mainnet would limit scalability of sharding.
    - EVM not well suited to managing sharding infrastructure efficiently (bit operations, etc.)
    - Recent innovations in crypto enable much more ambitious design.
  - Project reasons
    - Huge inertia to making changes on Mainnet.
    - Convergence of infrastructure for Proof of Stake and Sharding.
    - Separate chain allows radical changes and side-steps technical debt.

**Main Chain**  
provides staking

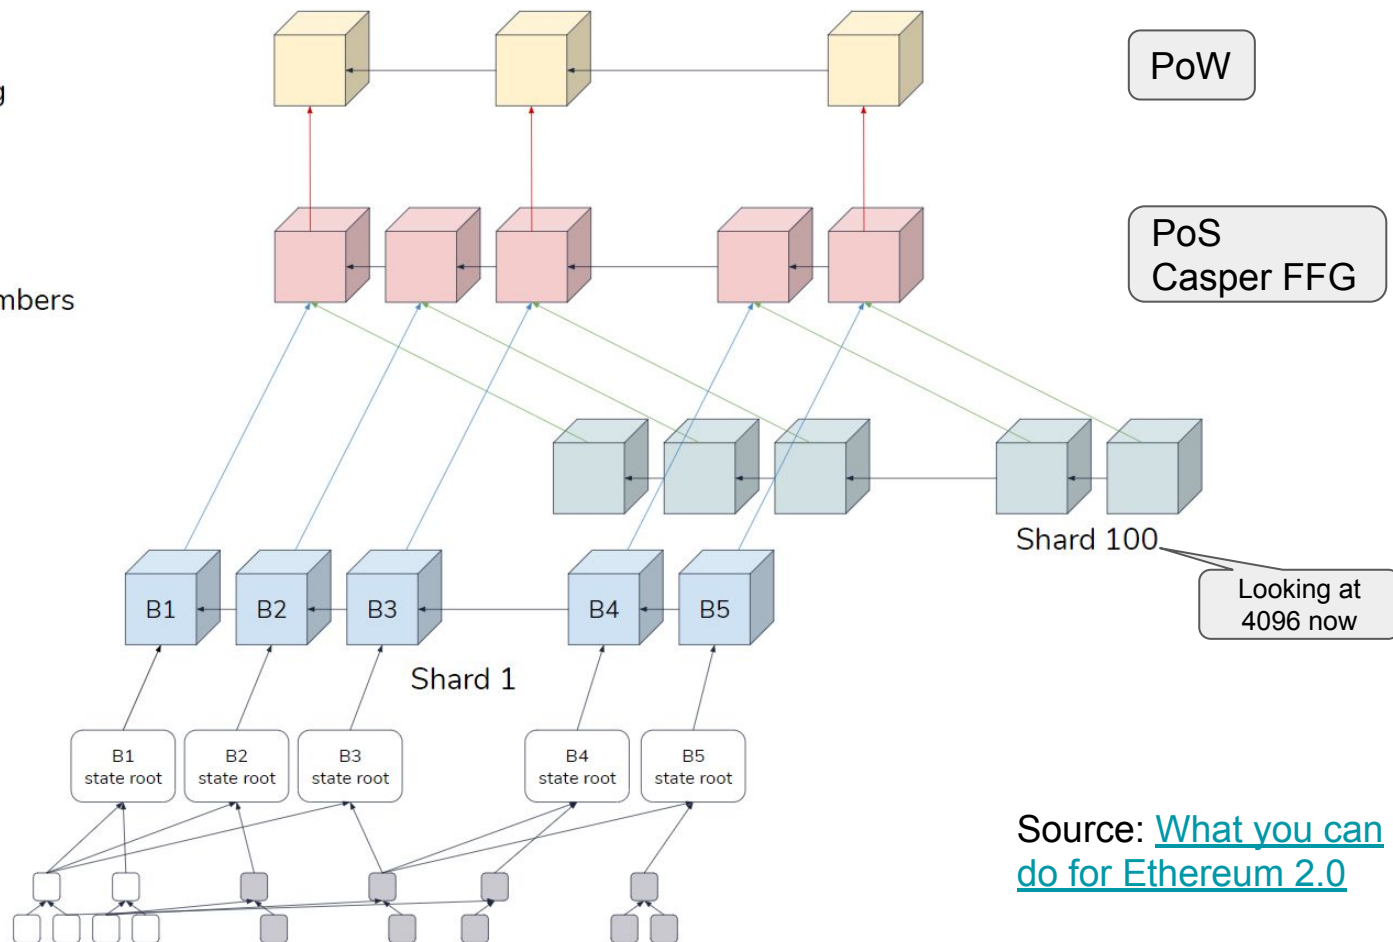
PoW

**Beacon Chain**  
provides random numbers

PoS  
Casper FFG

**Shard Chain**  
provides data

**VM**  
provides state  
execution result



Source: [What you can do for Ethereum 2.0](#)

# Differences and discontinuities 1.0 → 2.0

Changes are proposed to:

- Consensus algorithms
- Concurrency
- State transition
- Cryptoeconomics
- Cryptographic primitives
- P2P networking
- Contract programming model

(Implicit) Goals:

- Huge scalability
- Enhanced decentralisation
- Quantum security
- Protocol agility



# Consensus

Ethereum 1.0	Ethereum 2.0	Ethereum 2.x
<p>Proof of Work</p> <ul style="list-style-type: none"><li>• Huge environmental cost</li><li>• No finality</li><li>• Volatile block times</li><li>• Vulnerable to ASICs</li></ul>	<p>Proof of Stake overlay</p> <ul style="list-style-type: none"><li>• <a href="#">Beacon</a> and shard chains are PoS</li><li>• Periodic finality</li><li>• Casper FFG finality overlay on 1.0 chain</li><li>• Regular 5s block time</li></ul>	<p>Full <a href="#">Casper CBC</a> Proof of Stake.</p>

- Separating the beacon chain from Mainnet allows for vastly more validators.
  - Individual validator stake can be reduced from 1500 Eth to 32 Eth
  - 10% Eth staked => 300K validators
  - Aiming for ~4096 shards
- PoS validator and sharding validator pools can be combined, providing higher security levels.

# Concurrency

Ethereum 1.0	Ethereum 2.0	Ethereum 2.x
Globally single threaded <ul style="list-style-type: none"><li>Runs at speed <math>O(c)</math></li></ul>	Sharded state and parallel Tx processing <ul style="list-style-type: none"><li>Runs at speed <math>O(c^2)</math></li></ul>	Exponentially sharded?

- $O(c)$  is the processing power of a typical single laptop.
- Goal is to have no requirements that would prevent consumers on commodity hardware from participating in staking/validating. Promotes:
  - Network security
  - Decentralisation

# State transition

Ethereum 1.0	Ethereum 2.0	Ethereum 2.x
<p>The EVM</p> <ul style="list-style-type: none"><li>• State root recorded in every block.</li><li>• 256-bit architecture poor fit for most hardware.</li><li>• Limited compiler support.</li></ul>	<p><a href="#">eWASM</a></p> <ul style="list-style-type: none"><li>• Standardised, optimised.</li><li>• Wide range of implementations.</li><li>• Strong compiler support.</li><li>• Remove pre-compiles from the protocol.</li></ul>	<p>Alternative execution engines?</p>

Not yet decided: [delayed state execution](#).

- Blockchain only records the history of transactions.
- State can be calculated later on-demand (“lazy evaluation”).
- May allow for alternative execution engines on the public chain.



# Cryptoeconomics 1: slashing

Ethereum 1.0	Ethereum 2.0
Network security is largely incentive driven: <ul style="list-style-type: none"><li>• revolves around mining reward.</li></ul>	Network security assured by a combination of <ul style="list-style-type: none"><li>• Staking reward<ul style="list-style-type: none"><li>○ Rewards desirable behaviour</li></ul></li><li>• Slashing<ul style="list-style-type: none"><li>○ Punishes demonstrably bad behaviour</li></ul></li></ul>

- “Fraud proofs”/“fault proofs” become a thing. Participants are rewarded for reporting incorrect behaviour and perpetrators are slashed.

# Cryptoeconomics 2: storage rent

Ethereum 1.0	Ethereum 2.x
Blockchain storage is paid for once and persists forever.	<ul style="list-style-type: none"><li>• Charge ongoing fees for data and contract storage to reduce state size growth over time.</li><li>• Provide a way for users to re-upload locally stored state to the blockchain if it gets deleted.</li></ul>

- This is controversial and under much discussion.
- No other blockchain does this, but it seems inevitable in the long run.

# Crypto 1: aggregate block signatures

Ethereum 1.0	Ethereum 2.0	Ethereum 2.x
Not required	<a href="#">BLS aggregate signatures</a> allow for large numbers of validators to be handled efficiently. <ul style="list-style-type: none"><li>• Not quantum safe!</li></ul>	<a href="#">ZK-STARKs</a> <ul style="list-style-type: none"><li>• Quantum resistant</li></ul>

- Looking at the BLS12-381 elliptic curve for the aggregate signatures.
- [ZK-STARKs](#).
  - It is expected that STARKs will play a big role in the future in a number of protocol aspects.

# Crypto 2: randomness

Ethereum 1.0	Ethereum 2.0
<ul style="list-style-type: none"><li>• No in-protocol use of randomness.</li><li>• No easy way for applications to access randomness.</li></ul>	Protocol fundamentally relies on availability of a random beacon. Quality of randomness expected to be sufficient for many applications

Choice of implementations for the random beacon:

- BLS threshold signatures *a la* Dfinity
  - Complex set up (DKG)
  - Not quantum secure
- RANDAO (current proposal)
  - Needs a [VDF](#) (verifiable delay function) to avoid “last actor” issue. Blergh.

# Crypto 3: hash algorithm

Ethereum 1.0	Ethereum 2.0
Keccak256 (almost SHA3)	<u>MiMC?</u> <ul style="list-style-type: none"><li>• Minimises multiplications.</li><li>• Allows for more efficient STARK implementations.</li></ul>

Applications of hashing broaden in Ethereum 2.0:

- RANDAO hash onion for the beacon chain.
- Construction of proofs-of-custody of collation bodies.
- Provision of Merkle proofs in light and/or stateless clients.
- Quantum resistant transaction signatures?

# Crypto 4: account abstraction

Ethereum 1.0	Ethereum 2.x
Transaction validity is baked into the protocol: <ul style="list-style-type: none"><li>• Signature verification</li><li>• Nonce checking</li><li>• Fee verification</li></ul>	<a href="#">Account abstraction</a> : all accounts become contract accounts. Allow contracts to specify their own Tx validity criteria. E.g. <ul style="list-style-type: none"><li>• Alternative signature schemes</li><li>• Flexible fees payment options</li></ul>

This was [initially proposed](#) for the Mainnet in April 2016.

- As of the [last-but-one](#) Core Devs' it has now been iced until Ethereum 2.0
- An example of the difficulty of making disruptive changes to the current Mainnet.

# Networking

Ethereum 1.0	Ethereum 2.0
<i>devp2p</i> <ul style="list-style-type: none"><li>• creaky</li></ul>	Sharded p2p protocol likely based on <i>libp2p</i> . <ul style="list-style-type: none"><li>• Allow for fast shuffling of validators</li></ul>

# Programming model

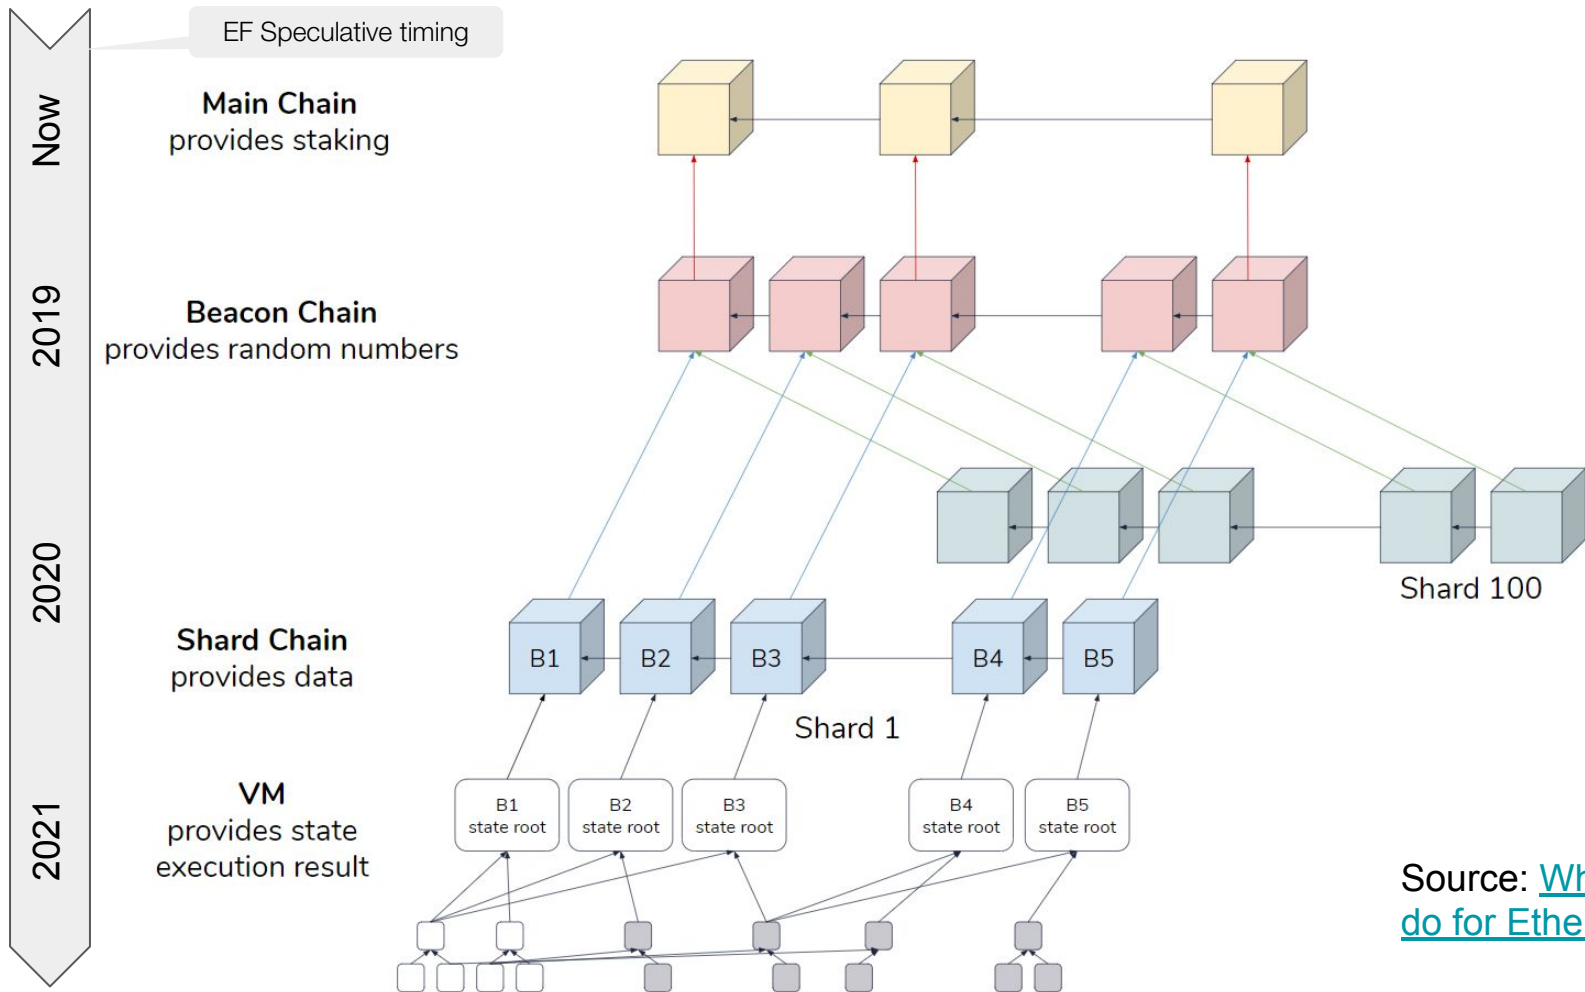
Ethereum 1.0	Ethereum 2.x
Inter-contract calls are synchronous.	<p>Undecided</p> <ul style="list-style-type: none"><li>• Sync within shards, async between?<ul style="list-style-type: none"><li>○ ✓ Natural to implement</li><li>○ ✗ Developers need to be shard-aware.</li></ul></li><li>• Async everywhere?<ul style="list-style-type: none"><li>○ ✓ Allows load-balance between shards.</li><li>○ ✗ A major break with the current model.</li></ul></li><li>• <u>Sync</u> everywhere?<ul style="list-style-type: none"><li>○ ✓ Maintains current programming model.</li><li>○ ✗ Complex to implement?</li></ul></li></ul>

In general, introducing asynchrony is likely to make auditing the correctness of contract systems substantially more difficult.



# Delivery Timeline

(Status: speculative. Based on [utterances](#) from the Ethereum Foundation)



Source: [What you can do for Ethereum 2.0](#)

## Phase 0: The PoS beacon chain (2019?)

- Implement PoS Beacon chain as a side chain
  - Random number generation: RANDAO using verifiable delay function.
  - Aggregate Signatures for validators
    - BLS aggregate signatures likely using BLS12-381 elliptic curve
  - Begin migration to new hash function(?)
    - Keccak256 -> MiMC hash?
  - 5s block time
  - New p2p network layer
- Brings Casper FFG periodic finality to Mainnet
- 32 Eth stake per validator
  - Slashing mechanism
- 80% Reduction in PoW mining reward (3 Eth→0.6 Eth?)

## Phase 1: Sharded transaction handling (The data layer - 2020?)

- Add sharding infrastructure for transactions
  - Each shard maintains an append-only log of transaction data.
  - Implement block proposer, block attester and block validator mechanisms.
  - Implement proofs of custody.
  - Implement cross links to beacon chain.
  - Enhance p2p networking to handle all of this.
- No state execution!
- Test the mechanism design.

## Phase 2: Sharded state handling (The execution layer - 2021?)

- Add state execution to shards
  - eWASM engine to replace EVM(?)
  - Delayed state execution(?)
  - Synchronous or asynchronous intra-shard contract calls?
  - Synchronous or asynchronous inter-shard contract calls?
  - Light clients and/or stateless clients.
- A fully usable,  $O(c^2)$  scaled infrastructure.
  - $O(c)$  is the compute power available to a consumer laptop or small VPS. The same speed as Mainnet today.

## Phase X: Also under discussion...

- ZK-STARKs.
  - Quantum resistant zero-knowledge proofs. Lots of use cases.
- Alternative execution engines.
- Storage rent.
- Full PoS: Casper CBC.
- Account abstraction.
- Mainnet becomes a side-chain to the beacon chain.
  - Abandon PoW?